

**ROB 534: Sequential Decision Making in Robotics, Winter 2025**  
**HW #1: Discrete and Sampling-based Planning**  
**Due: 1/29/25 (midnight)**

**Questions**

1. Search-based planning (20 pts)

(a) Consider the planning problem shown in Figure 1. Let '1' be the initial state, and let '6' be the goal state.

(i) By hand, use backward value iteration to determine the stationary cost-to-go (i.e, minimum cost between each state and the goal). **Show your work.**

(ii) Do the same but instead use forward value iteration. **Show your work.**





Figure 1: Six-state discrete planning problem.

(b) Consider the 6-DOF planning problem below. The goal is to get the end effector from a starting configuration (defined by five rotation angles  $R$  and one linear dimension  $P$ ) to an end configuration (also defined in the 6D space). The 6D space is discretized, and the cost metric is defined as:  $C = \Delta R + \Delta R_2 + \Delta R_3 + \Delta R_4 + \Delta R_5 + \Delta P$ , where  $\Delta R$  is the total angular distance travelled in the  $R$  degree of freedom,  $\Delta P$  is the total linear distance travelled, etc.



Which of the following heuristics is admissible for A\* search on this problem?

- i. Euclidean distance in 6DOF space from current node to goal
- ii. Euclidean distance in 6DOF space from current node to goal times 2
- iii. Euclidean distance in 6DOF space from current node to goal divided by 2
- v. Sum of total 6DOF distance travelled by optimal solution to goal state from current node

(c) For each heuristic that you deemed admissible in Part (b), order them from least informed to most informed. **Explain your answers.**

## Programming Assignment (Do not start at last minute!) (60 points)

The zip files on canvas contain helper functions and data for this assignment. Download them in addition to this word document. You may use any programming language you wish to complete the homework. Helper functions are provided in the zip file if you choose to use Matlab or Python. The text below is written describing the Matlab code; you should substitute the appropriate functions in Python if you use those functions.

In your Matlab program, nodes should be opaque data structures with operations `get_start.m`, `get_goal.m`, and `get_neighbors.m`. Restrict the number of nodes expanded to 10,000 to avoid very long run times.

To make your job easier, we have provided a Matlab priority queue implementation in the archive. You are free to use this implementation or not, as you see fit. The priority queue has the following functions:

- `pq init`: Initialize a priority queue.
- `pq set`: Reset the priority of an element or insert it if it's not already there.
- `pq pop`: Remove and return the first element. (The first element is the one with the smallest numerical priority value). It breaks ties arbitrarily.
- `pq test`: Check whether an item is already on the priority queue.

We have also provided two mazes (`maze1.pgm` and `maze2.pgm`) along with helper functions to read the mazes, convert from state  $(x,y)$  to index, plot paths, find neighbors and other operations. Familiarize yourself with the provided code before getting started.

**Step 1 (10 points): Before implementing any algorithms, draw a flowchart of the A\* and RRT algorithms.** These flowcharts should have text bubbles describing each step with arrows connecting them. Decisions should also be represented as branches with labels (e.g., yes/no). Include the flowcharts in your report.

### Step 2 (30 points):

i. **Implement A\* with an admissible heuristic. Which heuristic did you use? Provide an image of the path found by your A\* implementation on both mazes.** Have the starting node be the top left cell:  $(x, y)$  coordinates  $(1, 1)$ , and have the ending cell be the bottom right cell:  $(x,y)$  coordinates  $(\text{Cols}, \text{Rows})$ . Use the total distance traversed by the robot as the cost function.

ii. Allow A\* to multiply your admissible heuristic by a constant value epsilon (greedy A\*). **Set up a program that does the following:**

1. A\* runs with a user-provided epsilon
2. Once search is completed it then sets a  $\text{new\_epsilon} = \text{epsilon} - 0.5 * (\text{epsilon} - 1)$ . If  $\text{new\_epsilon}$  is less than 1.001, it sets  $\text{new\_epsilon} = 1$ .
3. A\* runs again with the new epsilon value

4. The loop continues with deflating epsilon until either (a) a running time limit is reached or (b) the search has completed with  $\epsilon = 1$ .

**Use a starting value of  $\epsilon = 10$ . Record for running time limits 0.05 seconds, 0.25 seconds, and 1 second (on both maps and for each epsilon value where the search completed): number of nodes expanded and path length.**

iii. **Implement the RRT algorithm to solve the 2D problem in continuous space on maze1.pgm and maze2.pgm [1].** You may use the knnsearch function in the statistics toolbox for nearest neighbors if using Matlab. Hint: you should sample the goal periodically to get goal directed behavior (do not use fully uniform sampling). You may assume the goal is reached if the x and y values are both within 1 of the goal. You may also assume the robot is modeled as a point. The provided check\_hit.m gives an obstacle test. **Provide an image of a path from the RRT algorithm on both maps. Also provide the path length and running time.**

**Step 3 (20 points):** For this problem, we will expand the state space to 4D to incorporate the dynamics of the agent moving through the world. The agent's state will now consist of (x,y) coordinates as well as (dx,dy) velocities. Instead of controlling its movement directly, the agent now controls its acceleration and deceleration. We have restricted the problem as follows:

- The agent starts in the top left corner  $(x,y) = (1,1)$  with zero velocity  $(dx, dy) = (0, 0)$ .
- The agent must move to the bottom right corner  $(x,y) = (\text{Rows}, \text{Cols})$  and slow down to zero velocity.
- The agent can accelerate or decelerate by one in either dx or dy (or remain at the same velocity).
- The agent's velocity in both directions must always remain less than or equal to a maximum velocity provided in the variable  $\text{maxV} = 2$ . Both velocity components must also always remain positive.
- The agent cannot slow down by hitting an obstacle. If it is traveling at a high speed, it must slow down several squares before an obstacle to avoid collision. (You do not want to scratch your car on those pesky obstacles).

For example, if the agent is in state  $(x,y,dx,dy) = (1,1,2,0)$ , it can choose to speed up dy, slow down dx, or remain at the same speed. It cannot speed up dx because it would break the speed limit, and it cannot slow down dy because it would be moving backwards. If it speeds up dy, it goes to  $(x, y, dx, dy) = (3,2,2,1)$ . If it slows down dx, it goes to  $(x,y,dx,dy) = (2,1,1,0)$ . If it remains at the same speed, it goes to  $(x, y, dx, dy) = (3, 1, 2, 0)$ . For more clarification, see the file get\_neighbors\_dynamic.m in the archive.

The 4D problem uses the same maps as before, but they are now read in using the read\_map\_for\_dynamics.m function. We have also provided you with the functions test\_dynamic\_neighbors.m and test\_dynamic\_path.m to help familiarize you with the format of the maps in 4D space.

For this problem, we have provided you with Matlab code for the functions `get_goal_dynamic.m`, `get_start_dynamic.m`, and `get_neighbors_dynamic.m`. These will replace the corresponding functions which you implemented for the 2D problem. The functions `dynamic_state_from_index.m` and `dynamic_index_from_state.m` also provide the same functionality as the corresponding 2D functions.

**i. Replace the A\* functions that you coded up for the 2D maze problem above and run your search algorithms on the 4D problem to minimize time (not distance). Develop an informed and admissible heuristic for the 4D domain (you only need one). What was your heuristic? Provide an image of your path on both mazes.**

**ii. Apply your deflating heuristic program to the 4D problem. Use a starting value of  $\epsilon = 10$ . Record for running time limits 0.05 seconds, 0.25 seconds, and 1 second (on both maps and for each epsilon value where the search completed): number of nodes expanded and path length.**

**Discussion Questions (20 points):** Address these questions (1-2 paragraphs each):

- 1) Consider a 4D grid problem with constant-time constraints (i.e., limits on time to compute). Informally describe the problem and solution using the constant-time motion planning (CTMP) algorithm [1]. Additionally, discuss potential real-world applications of CTMP beyond those highlighted in the research paper.
- 2) Consider using RRT-Connect on the 4D problem. What advantages would this have over standard RRT? What challenges would it lead to that would need to be overcome? Hint: Review the RRT\*-Connect paper [3].
- 3) What modifications to the A\* algorithm would you make if the robot discovered the environment as it went along (i.e., obstacles appeared and disappeared when the robot was near them)? You can assume the world itself is static, but the robot discovers the world as it moves. Do not provide full pseudocode, just a high-level description. Hint: Review the D\*-Lite algorithm [4].
- 4) What modifications to the RRT algorithm would you make if the robot's position were uncertain (partially observable)? Do not provide full pseudocode, just a high-level description. Hint: Review the RRBT algorithm [5].

## References

- [1] Islam, F., Salzman, O., Agarwal, A., & Likhachev, M., Provably constant-time planning and replanning for real-time grasping objects off a conveyor belt. *The International journal of robotics research*, 40(12-14), 1370-1384, 2021.
- [2] R. Mashayekhi, M.Y.I. Idris, M.H. Anisi, I. Ahmedy and I. Ali. "Informed RRT\*-connect: An asymptotically optimal single-query path planning method." *IEEE Access*, vol. 8, pp. 19842-19852, 2020.
- [3] S. Koenig and M. Likhachev, "D\* Lite," *Proc. AAAI/IAAI*, Vol. 15, 2002.
- [4] A. Bry, and N. Roy. "Rapidly-exploring random belief trees for motion planning under uncertainty." *Proc. IEEE Conference on Robotics and Automation (ICRA)*, 2011.

**Questions:**

1. Search-based planning (20 pts)
  - a. Consider the planning problem
    - i. Use backward value iteration to determine the stationary cost-to-go

Definitions:

$J(s)$  = cost-to-go for state  $s$

$C(s,s')$  = cost to transition from  $s$  to  $s'$

$\text{Succ}(s)$  = set of successor states for state  $s$

$J(s) = \min_{s' \in \text{Succ}(s)} \{C(s,s') + J(s')\}$

Transition Costs:

$C(1,2) = 3$ ,  $C(1,3) = 2$ ,  $C(1,5) = 1$ ,

$C(2,3) = 2$ ,  $C(2,4) = 1$ ,

$C(3,4) = 3$ ,  $C(3,5) = 4$ ,

$C(4,6) = 1$ ,

$C(5,6) = 5$ ,

$C(6,6) = 1$ .

Initialization:

$J(6) = 0$ , All other states are initialized with  $J(s) = \text{infinity}$

Iteration 1:

$J(6) = 0$

$J(5) = \min\{C(5,6) + J(6)\} = \min\{5 + 0\} = 5$

$J(4) = \min\{C(4,6) + J(6)\} = \min\{1 + 0\} = 1$

$J(3) = \min\{C(3,4) + J(4), C(3,5) + J(5)\} = \min\{3 + 1, 4 + 5\} = 4$

$J(2) = \min\{C(2,3) + J(3), C(2,4) + J(4)\} = \min\{2 + 4, 1 + 1\} = 2$

$J(1) = \min\{C(1,2) + J(2), C(1,3) + J(3), C(1,5) + J(5)\} = \min\{3 + 2, 2 + 4, 1 + 5\} = 5$

Iteration 2:

$J(5) = 5$  (no change)

$J(4) = 1$  (no change)

$J(3) = \min\{C(3,4) + J(4), C(3,5) + J(5)\} = \min\{3 + 1, 4 + 5\} = 4$  (no change)

$J(2) = \min\{C(2,3) + J(3), C(2,4) + J(4)\} = \min\{2 + 4, 1 + 1\} = 2$  (no change)

$J(1) = \min\{C(1,2) + J(2), C(1,3) + J(3), C(1,5) + J(5)\} = \min\{3 + 2, 2 + 4, 1 + 5\} = 5$   
(no change)

Final cost-to-go:

$J(1) = 5$ ,  $J(2) = 2$ ,  $J(3) = 4$ ,  $J(4) = 1$ ,  $J(5) = 5$ ,  $J(6) = 0$ .

ii. Do the same but with forward value iteration

Initialization:

$$J(1) = 0$$

All other states are initialized with  $J(s) = \text{infinity}$  except  $J(6)$  if explicitly updated during the iterations

Iteration 1:

$$J(1) = 0 \text{ (initial state)}$$

Compute costs to successors of state 1:

$$J(2) = \min\{C(1,2) + J(1)\} = \min\{3 + 0\} = 3$$

$$J(3) = \min\{C(1,3) + J(1)\} = \min\{2 + 0\} = 2$$

$$J(5) = \min\{C(1,5) + J(1)\} = \min\{1 + 0\} = 1$$

Iteration 2:

Update successors of state 2:

$$J(3) = \min\{J(3), C(2,3) + J(2)\} = \min\{2, 2 + 3\} = 2 \text{ (no change)}$$

$$J(4) = \min\{C(2,4) + J(2)\} = \min\{1 + 3\} = 4$$

Update successors of state 3:

$$J(4) = \min\{J(4), C(3,4) + J(3)\} = \min\{4, 3 + 2\} = 4 \text{ (no change)}$$

$$J(5) = \min\{J(5), C(3,5) + J(3)\} = \min\{1, 4 + 2\} = 1 \text{ (no change)}$$

Update successors of state 5:

$$J(6) = \min\{C(5,6) + J(5)\} = \min\{5 + 1\} = 6$$

Iteration 3:

Update successors of state 4:

$$J(6) = \min\{J(6), C(4,6) + J(4)\} = \min\{6, 1 + 4\} = 5$$

Update successors of state 6:

$$J(6) = \min\{J(6), C(6,6) + J(6)\} = \min\{5, 1 + 5\} = 5 \text{ (no change)}$$

Final cost-to-go:

$$J(1) = 0, J(2) = 3, J(3) = 2, J(4) = 4, J(5) = 1, J(6) = 5$$

b. Consider the 6-DOF planning problem - Admissibility Heuristic Evaluation:

- i. This heuristic is **admissible** because the Euclidean distance is always less than or equal to the actual cost because the actual cost has to consider path constraints, joint limits, and other factors that typically increase travel distance compared to a straight-line path.

- ii. This heuristic is **not admissible** because multiplying the Euclidean distance times 2 is an overestimation of the actual cost in most cases which violates the admissibility condition.  
This heuristic is the **least informed** because it underestimates the true cost significantly.
- iii. This heuristic is **admissible** because dividing the Euclidean distance by 2 results in an underestimation of the actual cost and since it is guaranteed to never overestimate the actual cost, it is admissible.
- iv. This heuristic is **admissible** because it exactly matches the actual cost because it assumes knowledge of the optimal path and since it never overestimates it is admissible.

**Summary: heuristics i, iii, iv are admissible and heuristic ii is not admissible**

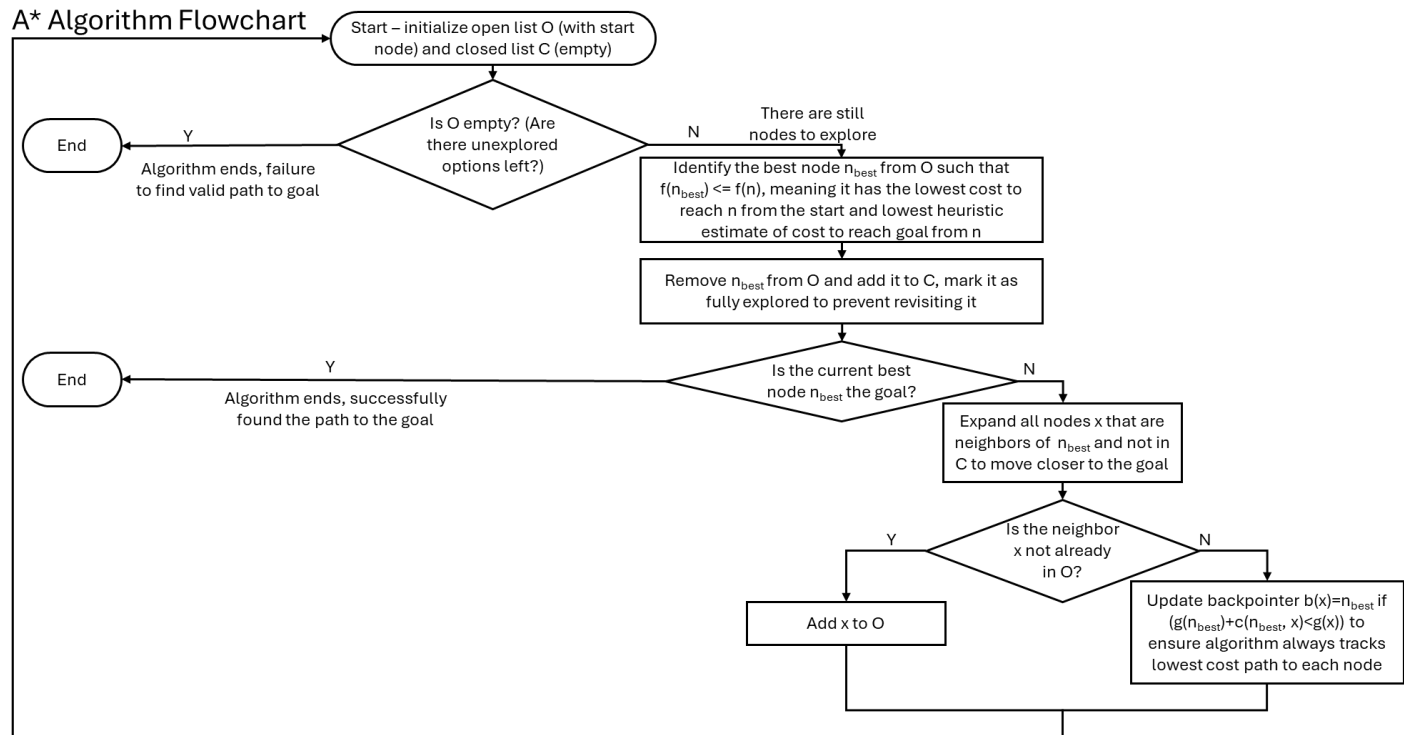
- c. Consider the 6-DOF planning problem - Least to Most Informed Heuristics:
  - i. This heuristic is **moderately informed**. It is more informed than heuristic iii because it can overestimate the actual cost but it can be equal to the actual cost so the lower bound is accurate since it directly represents the straight line distance to the goal.
  - ii. This heuristic is **more informed**. It is more informed than heuristics i and iii because it generally provides values closer to the actual cost since
  - iii. This heuristic is the **least informed**. It is the least informed since it significantly underestimates the actual cost and provides the weakest guidance.
  - iv. This heuristic is the **most informed** since it matches the true cost-to-go exactly in the optimal case.

**Summary: least to most informed:  $h(iii) < h(i) < h(ii) < h(iv)$**

## Programming Assignment:

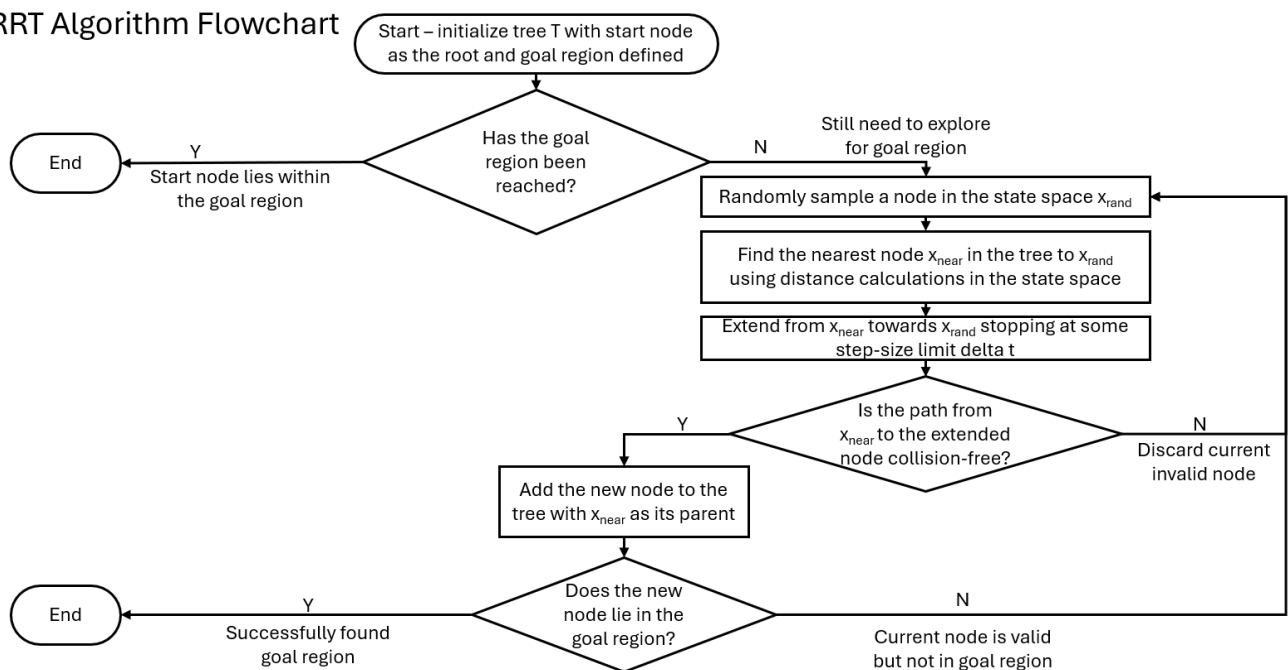
### Step 1: Flowchart of A\*:

#### A\* Algorithm Flowchart



### Flowchart of RRT:

#### RRT Algorithm Flowchart




### Step 2:


i. I used the Manhattan distance heuristic



## A\* Solution: Maze 1



## A\* Solution: Maze 2



ii. The number of nodes expanded was the same for all time limits for map 1 and map 2 for all values of epsilon.

At epsilon = 10 for time limits: 0.05 seconds, 0.25 seconds, and 1 second

- Map 1 number of nodes expanded: 128
- Map 1 path length: 57
- Map 2 number of nodes expanded: 115
- Map 2 path length: 47


At epsilon = 1 for time limits: 0.05 seconds, 0.25 seconds, and 1 second

- Map 1 number of nodes expanded: 409
- Map 1 path length: 47
- Map 2 number of nodes expanded: 389
- Map 2 path length: 47

```
Running with time limit 0.05s and initial epsilon 10
Epsilon: 10, Nodes expanded: 128, Path length: 57, Elapsed time: 0.0010
Running with time limit 0.25s and initial epsilon 10
Epsilon: 10, Nodes expanded: 128, Path length: 57, Elapsed time: 0.0010
Running with time limit 1s and initial epsilon 10
Epsilon: 10, Nodes expanded: 128, Path length: 57, Elapsed time: 0.0020
Testing Maze 2 with epsilon decay...
Running with time limit 0.05s and initial epsilon 10
Epsilon: 10, Nodes expanded: 115, Path length: 47, Elapsed time: 0.0000
Running with time limit 0.25s and initial epsilon 10
Epsilon: 10, Nodes expanded: 115, Path length: 47, Elapsed time: 0.0010
Running with time limit 1s and initial epsilon 10
Epsilon: 10, Nodes expanded: 115, Path length: 47, Elapsed time: 0.0000
Testing Maze 1 with epsilon decay...
Running with time limit 0.05s and initial epsilon 1
Epsilon: 1, Nodes expanded: 409, Path length: 47, Elapsed time: 0.0020
Running with time limit 0.25s and initial epsilon 1
Epsilon: 1, Nodes expanded: 409, Path length: 47, Elapsed time: 0.0020
Running with time limit 1s and initial epsilon 1
Epsilon: 1, Nodes expanded: 409, Path length: 47, Elapsed time: 0.0021
Testing Maze 2 with epsilon decay...
Running with time limit 0.05s and initial epsilon 1
Epsilon: 1, Nodes expanded: 389, Path length: 47, Elapsed time: 0.0020
Running with time limit 0.25s and initial epsilon 1
Epsilon: 1, Nodes expanded: 389, Path length: 47, Elapsed time: 0.0020
Running with time limit 1s and initial epsilon 1
Epsilon: 1, Nodes expanded: 389, Path length: 47, Elapsed time: 0.0020
```

This terminal output shows the values above.

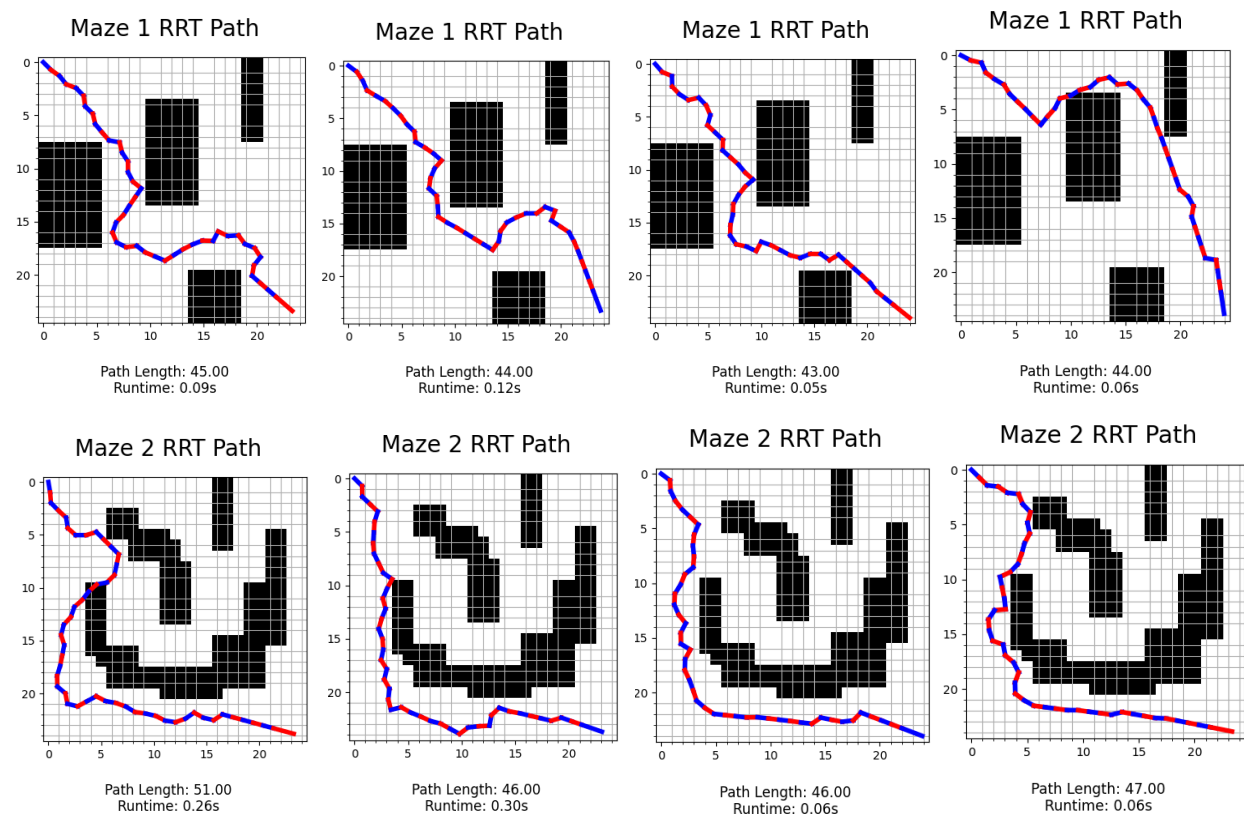
### iii. RRT algorithm



The terminal shows the path length and runtime with more significant figures.

```
PS C:\Users\chels\Downloads\HW1_complete\HW1_complete\hw1_python\hw1_python\provided_code> & c:/Users/chels/
Users/chels/Downloads/HW1_complete/HW1_complete/hw1_python/hw1_python/provided_code/test_maze_step2_iii.py
Running RRT with step size 1 and max iterations 1000 on Maze 1...
Iteration 0/1000, Tree size: 2
RRT Path Found for Maze 1. Path Length: 39.0, Runtime: 0.05655813217163086 seconds.
RRT Path Found for Maze 1. Path Length: 39.0, Runtime: 0.05655813217163086 seconds.
Running RRT with step size 1 and max iterations 1000 on Maze 2...
Iteration 0/1000, Tree size: 2
RRT Path Found for Maze 2. Path Length: 43.0, Runtime: 0.1022193431854248 seconds.
RRT Path Found for Maze 1. Path Length: 43.0, Runtime: 0.1022193431854248 seconds.
PS C:\Users\chels\Downloads\HW1_complete\HW1_complete\hw1_python\hw1_python\provided_code>
```


Below are other runs where the program only found a path for either map 1 or map 2




Step 3:

i. I used the Manhattan distance heuristic

A\* Solution: Maze 1 4D



A\* Solution: Maze 2 4D



ii. The number of nodes expanded was the same for all time limits for map 1 and map 2 for all values of epsilon.

At epsilon = 10 for time limits: 0.05 seconds, 0.25 seconds, and 1 second

- Map 1 number of nodes expanded: 237
- Map 1 path length: 51
- Map 2 number of nodes expanded: 1330
- Map 2 path length: 52

At epsilon = 1 for time limits: 0.05 seconds, 0.25 seconds, and 1 second

- Map 1 number of nodes expanded: 379
- Map 1 path length: 38
- Map 2 number of nodes expanded: 1883
- Map 2 path length: 36

The terminal output shows the values above:

```
Testing Maze 1 ...
Initial epsilon: 10, Nodes expanded: 237, Path length: 51, Elapsed time: 0.0025, Time limit: 0.05s
Initial epsilon: 10, Nodes expanded: 237, Path length: 51, Elapsed time: 0.0020, Time limit: 0.25s
Initial epsilon: 10, Nodes expanded: 237, Path length: 51, Elapsed time: 0.0010, Time limit: 1s
Initial epsilon: 1, Nodes expanded: 379, Path length: 38, Elapsed time: 0.0050, Time limit: 0.05s
Initial epsilon: 1, Nodes expanded: 379, Path length: 38, Elapsed time: 0.0055, Time limit: 0.25s
Initial epsilon: 1, Nodes expanded: 379, Path length: 38, Elapsed time: 0.0050, Time limit: 1s

Testing Maze 2 ...
Time limit 0.05s reached.
Initial epsilon: 10, Nodes expanded: 1330, Path length: 52, Elapsed time: 0.1308, Time limit: 0.05s
Initial epsilon: 10, Nodes expanded: 1330, Path length: 52, Elapsed time: 0.1300, Time limit: 0.25s
Initial epsilon: 10, Nodes expanded: 1330, Path length: 52, Elapsed time: 0.1309, Time limit: 1s
Time limit 0.05s reached.
Initial epsilon: 1, Nodes expanded: 1883, Path length: 36, Elapsed time: 0.0814, Time limit: 0.05s
Initial epsilon: 1, Nodes expanded: 1883, Path length: 36, Elapsed time: 0.0780, Time limit: 0.25s
Initial epsilon: 1, Nodes expanded: 1883, Path length: 36, Elapsed time: 0.0793, Time limit: 1s
```

Python file is attached at the end of this document!

### Discussion Questions:

1. Consider a 4D grid problem with constant-time constraints (i.e., limits on time to compute). Informally describe the problem and solution using the constant-time motion planning (CTMP) algorithm [1]. Additionally, discuss potential real-world applications of CTMP beyond those highlighted in the research paper.

For a 4D grid problem with constant-time constraints, the objective is to find a path for a robot to traverse through a 3D grid with limited time to compute each action. The 4D grid is represented by  $(x, y, dx, dy)$  where  $x$  and  $y$  are spatial coordinates and  $dx$  and  $dy$  are velocities. The problem is the time required for computing each action and the CTMP algorithm solves this problem by ensuring each action takes a constant amount of time to compute, regardless of how complex the problem is. It uses simplifications to allow for faster, more efficient processing which means that there is some trade off between the level of optimality of the solution to make it faster to provide more real-time feasibility.

The main real-world applications highlighted by the research paper are for moving objects off conveyor belts in warehouse and manufacturing applications. The CTMP algorithm can have several other real-world applications such as multi-agent systems, video games, simulations, navigation in unpredictable environments which is useful for domains such as agriculture, autonomous vehicles, human-robot-interaction, and military applications. For example, autonomous vehicles have to navigate through complex, unpredictable environments and make real-time decisions and the CTMP algorithm could help make these fast decisions where computation time is critical for safety and effectiveness. The same thing applies for military robots and drones that have to navigate unstructured, dynamic environments and make real-time decisions. Some agricultural robots don't function in real-time since they don't operate with people around, but for better human-robot interaction, safety, and effectiveness using the CTMP algorithm could help make faster decisions. For video games and real-time simulations the CTMP algorithm could help make decisions fast enough that things move fluidly in real-time. For multi-agent systems, fast computation is essential to maintain coordination and efficiency in unpredictable environments especially for mission-critical scenarios such as search and rescue operations.

2. Consider using RRT-Connect on the 4D problem. What advantages would this have over standard RRT? What challenges would it lead to that would need to be overcome? Hint: Review the RRT\*-Connect paper [3].

The RRT algorithm can be slow and sometimes does not find the goal, especially in 4D. The RRT-Connect algorithm offers several advantages and is more efficient. Because it grows two trees simultaneously from the start and goal states, it significantly reduces the search time and makes it easier to find a feasible path in fewer steps. The main advantages of RRT-Connect over RRT are faster convergence, more efficient exploration, lower computation cost, improved path quality, and better handling of large spaces.

The challenges with RRT-Connect are lack of optimality, inefficient use of cost metrics, suboptimal connection strategy, static exploration bias, and lack of guarantee of path quality. The RRT-Connect algorithm terminates as soon as it finds a path that connects the tree growing from the start to the tree growing from the goal which might not be the shortest or most optimal path. RRT\*-Connect improves on RRT-Connect and by using cost-based rewiring to refine paths based on a cost function which improves optimality and quality of the path and more efficiently uses cost metrics. RRT\*-Connect uses the rewiring to improve paths and results in an increasingly optimal solution as more samples are added.

3. What modifications to the A\* algorithm would you make if the robot discovered the environment as it went along (i.e., obstacles appeared and disappeared when the robot was near them)? You can assume the world itself is static, but the robot discovers the world as it moves. Do not provide full pseudocode, just a high-level description. Hint: Review the D\*-Lite algorithm [4].

Some modifications to the A\* algorithm that could help for dynamically exploring the environment are incremental path replanning, heuristic consistency, lazy evaluation, and path cost re-evaluation. Implementing incremental path planning allows the algorithm to dynamically update the path when new information from the environment is received such as obstacles appearing or disappearing as the robot drives towards or past obstacles. Ensuring a consistent heuristic that is admissible and monotonic during updates helps avoid invalid solutions and inefficiency. Lazy evaluation reuses previously generated path information to avoid recalculating paths from scratch as much as possible. Similarly, it could be modified to only recalculate affected nodes rather than the entire grid when the obstacle map changes and the path cost needs to be re-evaluated.

The D\*-Lite algorithm is specifically designed to handle these dynamic environments efficiently by using many of these modifications. It uses incremental path planning to minimize redundant calculations by reusing results from previous searches and only updates nodes affected by newly discovered

information using a priority queue to maintain and update these nodes. It uses dynamic cost updates for local edge costs when obstacles appear or disappear and efficiently propagates these changes to maintain path optimality. It uses a lazy evaluation strategy to prioritize nodes that are most likely to affect the current most optimal path rather than immediately exploring all possible nodes.

4. What modifications to the RRT algorithm would you make if the robot's position were uncertain (partially observable)? Do not provide full pseudocode, just a high-level description. Hint: Review the RRBT algorithm [5].

The main modifications that would need to be made are belief state planning instead of state space planning, uncertainty-aware sampling, propagation of belief nodes, probabilistic collision checks, and cost-to-go under uncertainty. When a robot's position is uncertain, the algorithm would need to be modified to not act in the state space and instead act in the belief space which represents a probability distribution over the possible positions the robot could be in which is estimated by using the robot's internal state estimates and sensor readings. Nodes in the RRT tree would represent the beliefs and to propagate belief nodes and expand the tree it would update the belief using techniques like Bayesian filtering with Kalman or particle filters. The algorithm would also need to do probabilistic collision checks to account for uncertainty in the robot's position and the representation of the environment. The cost function would also need to be modified to not only consider the path length but also the uncertainty at each state so it chooses paths that minimize uncertainty while also being efficient.

The RRBT algorithm addresses these modifications by operating in the belief space, using feedback controllers for stability, optimizing belief and cost, and re-evaluating belief nodes. The RRBT algorithm samples and expands in the belief space and constructs a tree of beliefs reflecting the uncertainty of the robot's state and the path/environment. It uses feedback controllers to reduce error propagation which ensures stability and reliability of planned paths even with the uncertainty. It reduces the trade off between path efficiency and uncertainty to minimize path length and positional error using an optimization framework that incorporates both cost and belief. When the RRBT algorithm receives new information, it re-evaluates belief nodes and edges to ensure they align with the updated beliefs, which continuously improves the solution quality.

## References:

- [1] Islam, F., Salzman, O., Agarwal, A., & Likhachev, M., Provably constant-time planning and replanning for real-time grasping objects off a conveyor belt. The International journal of robotics research, 40(12-14), 1370-1384, 2021.
- [2] R. Mashayekhi, M.Y.I. Idris, M.H. Anisi, I. Ahmedy and I. Ali. "Informed RRT\*-connect: An asymptotically optimal single-query path planning method." IEEE Access, vol. 8, pp. 19842-19852, 2020.
- [3] S. Koenig and M. Likhachev, "D\* Lite," Proc. AAAI/IAAI, Vol. 15, 2002.
- [4] A. Bry, and N. Roy. "Rapidly-exploring random belief trees for motion planning under uncertainty." Proc. IEEE Conference on Robotics and Automation (ICRA), 2011.
- [5] Geoff/Rakesh, there was a [5] but no reference with it, just letting you know, maybe the numbering is wrong because 2 doesn't appear in the in-text citations so maybe the in-text citations should be 1, 2, 3, 4 instead of 1, 3, 4, 5

## SDM\_HW1\_Chelse.py

```

1  #!/usr/bin/env python3
2
3  import abc
4  import numpy as np
5  import time
6  import heapq
7  import random
8  from sklearn.neighbors import NearestNeighbors
9  from matplotlib import pyplot as plt
10 from matplotlib.ticker import MultipleLocator
11 from heapq import heappop, heappush
12
13 class Maze(abc.ABC):
14     """ Base Maze Class """
15
16     def __init__(self, maze_array, start_index=None, goal_index=None):
17         """
18             maze_array - 2D numpy array with 1s representing free space
19                         0s representing occupied space
20         """
21         self.maze_array = maze_array
22         self.cols, self.rows = self.maze_array.shape
23         self.start_index = start_index
24         self.goal_index = goal_index
25
26     def __repr__(self):
27         if isinstance(self, Maze2D):
28             output = "2D Maze\n"
29             output += str(self.maze_array)
30             return output
31
32     @classmethod
33     def from_pgm(cls, filename):
34         """
35             Initializes the Maze from a (8 bit) PGM file
36         """
37         with open(filename, 'r', encoding='latin1') as infile:
38             header = infile.readline()
39             width, height, _ = [int(item) for item in header.split()[1:]]
40             image = np.fromfile(infile, dtype=np.uint8).reshape((height, width)) / 255
41
42             return cls(image.T)
43
44     def plot_maze(self):
45         """ Visualizes the maze """
46         self.plot_path([], "Maze")
47
48     def plot_path(self, path, title_name=None, runtime=None, path_length=None):

```



```

49     """
50     Plots the provided path on the maze and optionally shows
51     runtime and path length in a text box outside the plot area.
52     """
53     fig = plt.figure(1)
54     ax1 = fig.add_subplot(1, 1, 1)
55
56     spacing = 1.0 # Spacing between grid lines
57     minor_location = MultipleLocator(spacing)
58
59     # Set minor tick locations.
60     ax1.yaxis.set_minor_locator(minor_location)
61     ax1.xaxis.set_minor_locator(minor_location)
62
63     # Set grid to use minor tick locations.
64     ax1.grid(which='minor')
65
66     colors = ['b', 'r']
67     plt.imshow(self.maze_array.T, cmap=plt.get_cmap('bone'))
68
69     if title_name is not None:
70         fig.suptitle(title_name, fontsize=20)
71
72     # cast path to numpy array so indexing is nicer
73     path = np.array(path)
74     for i in range(len(path) - 1):
75         cidx = i % 2
76         plt.plot([path[i, 0], path[i + 1, 0]], [path[i, 1], path[i + 1, 1]],
77 color=colors[cidx], linewidth=4)
78
79     # If runtime and path length are provided, add them outside the plot
80     if runtime is not None and path_length is not None:
81         text = f"Path Length: {path_length:.2f}\nRuntime: {runtime:.2f}s"
82
83         # Adjust position and move the text further down
84         fig.subplots_adjust(bottom=0.2) # Make more room for the text box below the graph
85         plt.figtext(0.5, 0.02, text, ha="center", fontsize=12) # Move the text further
86         down
87
88     plt.show()
89
90     def check_occupancy(self, state):
91         """ Returns True if there is an obstacle at state """
92         return self.maze_array[int(state[0]), int(state[1])] == 0
93
94     def get_goal(self):
95         """ Returns the index of the goal """
96         return self.goal_index
97
98     def get_start(self):
99         """ Returns the index of the start state """

```

```

97         return self.start_index
98
99     def check_hit(self, start, deltas):
100         """
101         Returns True if there are any occupied states between:
102         start[0] to start[0]+dx and start[1] to start[1]+dy
103         """
104         x, y = start
105         dx, dy = deltas
106
107         if (x < 0) or (y < 0) or (x >= self.cols) or (y >= self.rows):
108             return True
109
110         if self.maze_array[int(round(start[0])), int(round(start[1]))] == 0:
111             return True
112
113         if dx == 0.0 and dy == 0.0: # no actual movement
114             return False
115
116         norm = max(abs(dx), abs(dy))
117         dx /= norm
118         dy /= norm
119
120         for _ in range(int(norm)):
121             x += dx
122             y += dy
123             if (x < 0) or (y < 0) or (x >= self.cols) or (y >= self.rows):
124                 return True
125             if self.maze_array[int(x), int(y)] == 0:
126                 return True
127         return False
128
129     def check_occupancy(self, state):
130         """ Returns True if there is an obstacle at state """
131         return self.maze_array[int(state[0]), int(state[1])] == 0
132
133
134     class Maze2D(Maze):
135         """ Maze2D Class """
136
137         def __init__(self, maze_array, start_state=None, goal_state=None):
138             super().__init__(maze_array, start_state, goal_state)
139
140             if start_state is None:
141                 start_state = (0, 0)
142             self.start_state = start_state
143             self.start_index = self.index_from_state(self.start_state)
144
145             if goal_state is None:
146                 goal_state = (self.cols-1, self.rows-1)

```

```

147         self.goal_state = goal_state
148         self.goal_index = self.index_from_state(self.goal_state)
149
150     def index_from_state(self, state):
151         """ Gets a unique index for the state """
152         return state[0] * self.rows + state[1]
153
154     def state_from_index(self, state_id):
155         """ Returns the state at a given index """
156         x = int(np.floor(state_id / self.rows))
157         y = state_id % self.rows
158         return (x, y)
159
160     def get_neighbors(self, state_id):
161         """ Returns a List of indices corresponding to neighbors of a given state """
162         state = self.state_from_index(state_id)
163         deltas = [[0, -1], [0, 1], [-1, 0], [1, 0]]
164         neighbors = []
165         for delta in deltas:
166             if not self.check_hit(state, delta):
167                 new_state = (state[0] + delta[0], state[1] + delta[1])
168                 neighbors.append(self.index_from_state(new_state))
169         return neighbors
170
171     # Step 2 part i code
172     def a_star_step2_i(self):
173         """ Perform A* search to find the optimal path from start to goal """
174         start = self.start_index
175         goal = self.goal_index
176
177         open_set = []
178         heappush(open_set, (0, start)) # (f_score, state_index)
179
180         came_from = {}
181         g_score = {start: 0}
182         f_score = {start: self._heuristic(start)}
183
184         while open_set:
185             _, current = heappop(open_set)
186
187             if current == goal:
188                 return self._reconstruct_path(came_from, current)
189
190             for neighbor in self.get_neighbors(current):
191                 tentative_g_score = g_score[current] + 1
192
193                 if neighbor not in g_score or tentative_g_score < g_score[neighbor]:
194                     came_from[neighbor] = current
195                     g_score[neighbor] = tentative_g_score
196                     f_score[neighbor] = g_score[neighbor] + self._heuristic(neighbor)

```

```

197         heappush(open_set, (f_score[neighbor], neighbor))
198
199     return []
200
201     def _heuristic(self, state_index):
202         """ Calculate Manhattan distance heuristic to goal """
203         state = self.state_from_index(state_index)
204         goal = self.state_from_index(self.goal_index)
205         return abs(goal[0] - state[0]) + abs(goal[1] - state[1])
206
207     def _reconstruct_path(self, came_from, current):
208         """ Reconstruct path from start to goal """
209         path = [self.state_from_index(current)]
210         while current in came_from:
211             current = came_from[current]
212             path.append(self.state_from_index(current))
213         return path[::-1]
214
215     # Step 2 part ii code
216     def a_star_step2_ii(self, epsilon=1):
217         """ Perform A* search with a greedy heuristic adjusted by epsilon """
218         start, goal = self.start_index, self.goal_index
219
220         open_set = [(0, start)]
221         came_from, g_score, f_score = {}, {start: 0}, {start: self._greedy_heuristic(start,
epsilon)}
222
223         while open_set:
224             _, current = heappop(open_set)
225
226             if current == goal:
227                 return self._reconstruct_path(came_from, current), len(came_from)
228
229             for neighbor in self.get_neighbors(current):
230                 tentative_g_score = g_score[current] + 1
231
232                 if neighbor not in g_score or tentative_g_score < g_score[neighbor]:
233                     came_from[neighbor] = current
234                     g_score[neighbor] = tentative_g_score
235                     f_score[neighbor] = g_score[neighbor] + self._greedy_heuristic(neighbor,
epsilon)
236
237                     heappush(open_set, (f_score[neighbor], neighbor))
238
239             return [], 0 # No path found
240
241     def _greedy_heuristic(self, state_index, epsilon):
242         """ Calculate the greedy heuristic to the goal """
243         state = self.state_from_index(state_index)
244         goal = self.state_from_index(self.goal_index)
245         return epsilon * (abs(goal[0] - state[0]) + abs(goal[1] - state[1]))

```

```

245
246 def run_with_epsilon_decay(self, initial_epsilon=10, time_limits=[0.05, 0.25, 1]):
247     for time_limit in time_limits:
248         epsilon = initial_epsilon
249         start_time = time.time()
250         print(f"Running with time limit {time_limit}s and initial epsilon {epsilon}")
251
252         while time.time() - start_time < time_limit:
253             path, expanded_nodes = self.a_star_step2_ii(epsilon)
254             if path: # If path found
255                 break
256             # Decay epsilon
257             epsilon = max(1, epsilon - 0.5 * (epsilon - 1))
258
259             print(f"Epsilon: {epsilon}, Nodes expanded: {expanded_nodes}, Path length:
{len(path)}")
260
261 # Step 2 part iii code
262 def rrt(self, max_iter=1000, step_size=1):
263     def sample_free():
264         while True:
265             if random.random() < 0.2:
266                 return np.array(self.goal_state)
267             else:
268                 x = random.randint(0, self.cols - 1)
269                 y = random.randint(0, self.rows - 1)
270                 if self.check_occupancy((x, y)) == 0:
271                     return np.array([x, y])
272
273     def nearest_node(tree, sample):
274         nbrs = NearestNeighbors(n_neighbors=1, algorithm='ball_tree').fit(tree)
275         _, idx = nbrs.kneighbors([sample])
276         return tree[idx[0][0]]
277
278     def steer(from_node, to_node, step_size):
279         vector = to_node - from_node
280         distance = np.linalg.norm(vector)
281         if distance <= step_size:
282             return to_node
283         return from_node + vector / distance * step_size
284
285     # Initialize tree and time tracking
286     tree = np.array([self.start_state])
287     parent_map = {}
288     start_time = time.time()
289
290     for i in range(max_iter):
291         sample = sample_free()
292         nearest = nearest_node(tree, sample)
293         new_node = steer(nearest, sample, step_size)

```

```

294
295         if self.check_occupancy(nearest):
296             continue # Skip if new node hits an obstacle
297
298         tree = np.vstack([tree, new_node])
299         parent_map[tuple(new_node)] = tuple(nearest)
300
301         if np.linalg.norm(new_node - self.goal_state) <= 1:
302             path = [tuple(new_node)]
303             while tuple(nearest) != tuple(self.start_state):
304                 nearest = parent_map[tuple(nearest)]
305                 path.append(tuple(nearest))
306             path.reverse()
307
308             # Calculate path length
309             path_length = sum(np.linalg.norm(np.array(path[i+1]) - np.array(path[i]))
for i in range(len(path) - 1))
310             runtime = time.time() - start_time # Compute runtime
311             return path, path_length, runtime # Return the correct number of values
312
313         if i % 100 == 0:
314             print(f"Iteration {i}/{max_iter}, Tree size: {len(tree)}")
315
316         # Return an empty path if no valid path is found
317         runtime = time.time() - start_time # Compute runtime after all iterations
318         return [], 0, runtime # Return the correct number of values
319
320 # Step 3 code
321 class Maze4D(Maze):
322     """ Maze4D Class """
323
324     def __init__(self, maze_array, start_state=None, goal_state=None, max_vel=2):
325         super().__init__(maze_array, start_state, goal_state)
326
327         self.max_vel = max_vel
328
329         if start_state is None:
330             start_state = np.array((0, 0, 0, 0))
331         self.start_state = start_state
332         self.start_index = self.index_from_state(self.start_state)
333
334         if goal_state is None:
335             goal_state = np.array((self.cols - 1, self.rows - 1, 0, 0))
336         self.goal_state = goal_state
337         self.goal_index = self.index_from_state(self.goal_state)
338
339     def index_from_state(self, state):
340         """ Gets a unique index for the state """
341         velocities = self.max_vel + 1
342         return state[3] * self.rows * self.cols * velocities + \

```

```

343         state[2] * self.rows * self.cols + \
344         state[0] * self.rows + \
345         state[1]
346
347     def state_from_index(self, state_id):
348         """ Returns the state at a given index """
349         velocities = self.max_vel + 1
350         idx = state_id
351         dy = int(np.floor(idx / (self.rows * self.cols * velocities)))
352         idx -= dy * self.rows * self.cols * velocities
353         dx = int(np.floor(idx / (self.rows * self.cols)))
354         idx -= dx * self.rows * self.cols
355         x = int(np.floor(idx / self.rows))
356         y = idx % self.rows
357         return (x, y, dx, dy)
358
359     def get_neighbors(self, state_id):
360         """ Returns a List of indices corresponding to neighbors of a given state in the 4D
361         maze """
362         state = self.state_from_index(state_id)
363         neighbors = []
364
365         deltas = [
366             [-1, 0], [1, 0], # horizontal
367             [0, -1], [0, 1], # vertical
368             [-1, -1], [1, 1], # diagonal
369             [0, 0] # no movement
370         ]
371
372         for delta in deltas:
373             new_dx = state[2] + delta[0]
374             new_dy = state[3] + delta[1]
375
376             if new_dx < 0 or new_dx > self.max_vel or new_dy < 0 or new_dy > self.max_vel:
377                 continue
378
379             if not self.check_hit(state[0:2], [new_dx, new_dy]):
380                 new_state = (state[0] + new_dx, state[1] + new_dy, new_dx, new_dy)
381                 neighbors.append(self.index_from_state(new_state))
382
383         return neighbors
384
385     @staticmethod
386     def a_star_search_step3_i(maze, start_state, goal_state):
387         """ A* search for solving the maze """
388         def heuristic(state, goal_state):
389             return abs(state[0] - goal_state[0]) + abs(state[1] - goal_state[1])
390
391         def reconstruct_path(came_from, current_state):
392             path = [current_state]

```

```

392         while current_state in came_from:
393             current_state = came_from[current_state]
394             path.insert(0, current_state)
395         return path
396
397     open_set = []
398     heapq.heappush(open_set, (0, tuple(start_state)))
399     came_from = {}
400     g_score = {tuple(start_state): 0}
401     f_score = {tuple(start_state): heuristic(start_state, goal_state)}
402
403     while open_set:
404         _, current_state = heapq.heappop(open_set)
405
406         if current_state == tuple(goal_state):
407             return reconstruct_path(came_from, current_state)
408
409         for neighbor in maze.get_neighbors(maze.index_from_state(current_state)):
410             neighbor_state = maze.state_from_index(neighbor)
411             tentative_g_score = g_score[tuple(current_state)] + 1
412
413             if tuple(neighbor_state) not in g_score or tentative_g_score <
g_score[tuple(neighbor_state)]:
414                 came_from[tuple(neighbor_state)] = current_state
415                 g_score[tuple(neighbor_state)] = tentative_g_score
416                 f_score[tuple(neighbor_state)] = tentative_g_score +
heuristic(neighbor_state, goal_state)
417                 heapq.heappush(open_set, (f_score[tuple(neighbor_state)],
tuple(neighbor_state)))
418
419     return None # No path found
420
421     def a_star_step3_ii(self, epsilon=1):
422         """
423         Perform A* search with a greedy heuristic adjusted by epsilon.
424         The cost is based on the heuristic scaled by epsilon for the greedy part.
425         """
426         start = self.start_index
427         goal = self.goal_index
428
429         # Priority queue (min-heap)
430         open_set = []
431         heappush(open_set, (0, start)) # (f_score, state_index)
432
433         # Dictionaries for tracking costs and paths
434         came_from = {}
435         g_score = {start: 0}
436         f_score = {start: self._greedy_heuristic(start, epsilon)}
437
438         while open_set:
439             _, current = heappop(open_set)

```



```

440
441     # If we reach the goal, reconstruct the path
442     if current == goal:
443         return self._reconstruct_path(came_from, current), len(came_from)
444
445     # Process neighbors
446     for neighbor in self.get_neighbors(current):
447         tentative_g_score = g_score[current] + 1 # Assume uniform cost for grid steps
448
449         if neighbor not in g_score or tentative_g_score < g_score[neighbor]:
450             # Update cost to reach neighbor
451             came_from[neighbor] = current
452             g_score[neighbor] = tentative_g_score
453             f_score[neighbor] = g_score[neighbor] + self._greedy_heuristic(neighbor,
epsilon)
454
455             # Push to priority queue
456             heappush(open_set, (f_score[neighbor], neighbor))
457
458     return [], 0 # If no path is found
459
460 def _greedy_heuristic(self, state_index, epsilon):
461     """ Calculate the greedy heuristic to the goal, scaled by epsilon """
462     state = self.state_from_index(state_index)
463     goal = self.state_from_index(self.goal_index)
464     # Manhattan distance heuristic for 4D space, can adjust it for more complex
calculations
465     return epsilon * (abs(goal[0] - state[0]) + abs(goal[1] - state[1]) + abs(goal[2] -
state[2]) + abs(goal[3] - state[3]))
466
467 def _reconstruct_path(self, came_from, current):
468     """ Reconstruct path from start to goal """
469     path = [self.state_from_index(current)]
470     while current in came_from:
471         current = came_from[current]
472         path.append(self.state_from_index(current))
473     return path[::-1] # Reverse the path
474
475 def run_with_epsilon_decay(self, initial_epsilon=10, time_limits=[0.05, 0.25, 1]):
476     for time_limit in time_limits:
477         epsilon = initial_epsilon
478         start_time = time.time()
479         #print(f"Running with time limit {time_limit}s and initial epsilon {epsilon}")
480
481         while True:
482             #print(f"Running A* with epsilon={epsilon}")
483             path, expanded_nodes = self.a_star_step3_ii(epsilon)
484
485             elapsed_time = time.time() - start_time
486
487             # Check for time limit

```

```

488         if elapsed_time >= time_limit:
489             print(f"Time limit {time_limit}s reached.")
490             break
491
492         if path: # If a path was found
493             #print(f"Found path with {expanded_nodes} nodes expanded.")
494             #print(f"Path length: {len(path)}")
495             break
496
497         # Decay epsilon
498         new_epsilon = epsilon - 0.5 * (epsilon - 1)
499         if new_epsilon < 1.001:
500             new_epsilon = 1
501         epsilon = new_epsilon
502
503         print(f"Initial epsilon: {epsilon}, Nodes expanded: {expanded_nodes}, Path length:
504         {len(path)}, Elapsed time: {elapsed_time:.4f}, Time limit: {time_limit}s")
505
506 if __name__ == "__main__":
507     mazes = [
508         (Maze2D.from_pgm('maze1.pgm'), "Maze 1"),
509         (Maze2D.from_pgm('maze2.pgm'), "Maze 2"),
510     ]
511
512     # Setup and preprocessing for each maze
513     for maze, maze_name in mazes:
514         maze.start_state = (1, 1)
515         maze.goal_state = (maze.cols - 1, maze.rows - 1)
516         maze.start_index = maze.index_from_state(maze.start_state)
517         maze.goal_index = maze.index_from_state(maze.goal_state)
518
519     # Step 2 part i: Run A* on each maze
520     for maze, maze_name in mazes:
521         path = maze.a_star_step2_i()
522         print(f"Found path for {maze_name}")
523         maze.plot_path(path, f"A* Solution: {maze_name}")
524
525     # Step 2 part ii: Run with epsilon decay on each maze
526     for maze, maze_name in mazes:
527         print(f"Testing {maze_name} with epsilon decay...")
528         maze.run_with_epsilon_decay(initial_epsilon=10, time_limits=[0.05, 0.25, 1])
529
530     # Step 2 part iii: Run RRT on each maze
531     for maze, maze_name in mazes:
532         print(f"Running RRT with step size 1 and max iterations 1000 on {maze_name}...")
533         path, length, runtime = maze.rrt(max_iter=1000, step_size=1)
534         if path:
535             print(f"RRT Path Found for {maze_name}. Path Length: {length}, Runtime: {runtime}
seconds.")
536         maze.plot_path(path, f"{maze_name} RRT Path", runtime=runtime, path_length=length)

```

```
536         else:
537             print(f"RRT failed to find a path for {maze_name}")
538
539     # Step 3
540     mazes_4D = [
541         (Maze4D.from_pgm('maze1.pgm'), "Maze 1 4D"),
542         (Maze4D.from_pgm('maze2.pgm'), "Maze 2 4D")
543     ]
544     # Setup and preprocessing for each 4D maze
545     for maze, maze_name in mazes_4D:
546         # Test index_from_state and state_from_index for consistency
547         for x in range(maze.cols):
548             for y in range(maze.rows):
549                 for dx in range(3):
550                     for dy in range(3):
551                         state = (x, y, dx, dy)
552                         assert maze.state_from_index(maze.index_from_state(state)) == state,
553                             f"Mapping incorrect for state: {state}"
554
555     for maze, maze_name in mazes_4D:
556         # Step 3 part i: Solve using A* search
557         path = Maze4D.a_star_search_step3_i(maze, maze.start_state, maze.goal_state)
558         if path:
559             maze.plot_path(path, f'A* Solution: {maze_name}')
560         else:
561             print("No path found!")
562
563     for maze, maze_name in mazes_4D:
564         # Step 3 part ii: Run with epsilon decay on each 4D maze
565         print(f"Testing {maze_name} with epsilon decay...")
566         maze.run_with_epsilon_decay(initial_epsilon=1, time_limits=[0.05, 0.25, 1])
567
```